

Customer Segmentation and Profiling

December 23, 2021

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from scipy import stats
from sklearn.preprocessing import StandardScaler
from sklearn.manifold import TSNE
from sklearn.cluster import KMeans

from feature_engine.outlier_removers import Winsorizer
```

In C:\Users\olale\Anaconda3\lib\site-packages\matplotlib\mpl-data\stylelib_classic_test.mplstyle:

The text.latex.preview rcparam was deprecated in Matplotlib 3.3 and will be removed two minor releases later.

In C:\Users\olale\Anaconda3\lib\site-packages\matplotlib\mpl-data\stylelib_classic_test.mplstyle:

The mathtext.fallback_to_cm rcparam was deprecated in Matplotlib 3.3 and will be removed two minor releases later.

In C:\Users\olale\Anaconda3\lib\site-packages\matplotlib\mpl-data\stylelib_classic_test.mplstyle: Support for setting the 'mathtext.fallback_to_cm' rcParam is deprecated since 3.3 and will be removed two minor releases later; use 'mathtext.fallback : 'cm' instead.

In C:\Users\olale\Anaconda3\lib\site-packages\matplotlib\mpl-data\stylelib_classic_test.mplstyle:

The validate_bool_maybe_none function was deprecated in Matplotlib 3.3 and will be removed two minor releases later.

In C:\Users\olale\Anaconda3\lib\site-packages\matplotlib\mpl-data\stylelib_classic_test.mplstyle:

The savefig.jpeg_quality rcparam was deprecated in Matplotlib 3.3 and will be removed two minor releases later.

In C:\Users\olale\Anaconda3\lib\site-packages\matplotlib\mpl-data\stylelib_classic_test.mplstyle:

The keymap.all_axes rcparam was deprecated in Matplotlib 3.3 and will be removed two minor releases later.

In C:\Users\olale\Anaconda3\lib\site-packages\matplotlib\mpl-data\stylelib_classic_test.mplstyle:

The `animation.avconv_path` rcparam was deprecated in Matplotlib 3.3 and will be removed two minor releases later.

In `C:\Users\olale\Anaconda3\lib\site-packages\matplotlib\mpl-data\stylelib_classic_test.mplstyle`:

The `animation.avconv_args` rcparam was deprecated in Matplotlib 3.3 and will be removed two minor releases later.

```
[2]: df = pd.read_excel('Trifactor final dataset.xlsx')
df
```

```
[2]:
```

	Index	Date	event_type	product_id	category_id	\
0	0.0	2019-10-11	view	5016.0	183.0	
1	1.0	2019-10-27	view	4240.0	43.0	
2	2.0	2019-11-25	view	7721.0	208.0	
3	3.0	2019-11-12	view	597.0	71.0	
4	4.0	2019-11-13	view	13709.0	70.0	
...	
235996	235996.0	2019-11-17	cart	719.0	71.0	
235997	235997.0	2019-11-17	cart	718.0	71.0	
235998	235998.0	2019-11-16	cart	415.0	71.0	
235999	235999.0	2019-11-15	cart	4649.0	175.0	
236000	NaN	NaT	NaN	NaN	NaN	

		category_code	brand	price	user_id	\
0		appliances.kitchen.oven	artel	36.01	1095536.0	
1		electronics.video.tv	lg	2445.08	1153084.0	
2		appliances.environment.vacuum	philips	257.38	364298.0	
3		electronics.smartphone	huawei	163.20	3536496.0	
4		electronics.telephone	nokia	21.85	3542250.0	
...		
235996		electronics.smartphone	samsung	298.33	450109.0	
235997		electronics.smartphone	samsung	298.07	1516528.0	
235998		electronics.smartphone	samsung	94.96	2816015.0	
235999		appliances.kitchen.hood	bosch	144.15	247855.0	
236000		NaN	NaN	NaN	NaN	

	user_session	State	User_Score	Year	Quarter	Month	\
0	5342766.0	PA	3.0	2019.0	4.0	October	
1	4600705.0	CT	4.0	2019.0	4.0	October	
2	5991663.0	HI	3.0	2019.0	4.0	November	
3	9397681.0	NH	3.0	2019.0	4.0	November	
4	8711820.0	NV	1.0	2019.0	4.0	November	
...	
235996	9681143.0	NaN	NaN	2019.0	4.0	November	
235997	3956291.0	NaN	NaN	2019.0	4.0	November	
235998	12343925.0	NaN	NaN	2019.0	4.0	November	
235999	9159758.0	NaN	NaN	2019.0	4.0	November	

236000	NaN	NaN	NaN	NaN	NaN	NaN
--------	-----	-----	-----	-----	-----	-----

	Week of Year	Name of Day	Hour	Event_Time
0	41.0	Friday	16.0	16:57:06
1	44.0	Sunday	14.0	14:23:53
2	48.0	Monday	12.0	12:57:47
3	46.0	Tuesday	19.0	19:50:55
4	46.0	Wednesday	4.0	04:59:13
...
235996	47.0	Sunday	16.0	16:07:29
235997	47.0	Sunday	12.0	12:39:00
235998	46.0	Saturday	14.0	14:02:53
235999	46.0	Friday	9.0	09:17:27
236000	NaN	NaN	NaN	None

[236001 rows x 19 columns]

```
[3]: df.dtypes
```

```
[3]: Index          float64
Date            datetime64[ns]
event_type      object
product_id      float64
category_id     float64
category_code   object
brand           object
price           float64
user_id         float64
user_session    float64
State           object
User_Score      float64
Year            float64
Quarter         float64
Month           object
Week of Year    float64
Name of Day     object
Hour            float64
Event_Time      object
dtype: object
```

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 236001 entries, 0 to 236000
Data columns (total 19 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Index           236000 non-null float64
```

```

1  Date                236000 non-null  datetime64[ns]
2  event_type          236000 non-null  object
3  product_id          236000 non-null  float64
4  category_id         236000 non-null  float64
5  category_code       236000 non-null  object
6  brand               236000 non-null  object
7  price               236000 non-null  float64
8  user_id             236000 non-null  float64
9  user_session        236000 non-null  float64
10 State               204630 non-null  object
11 User_Score          204630 non-null  float64
12 Year                236000 non-null  float64
13 Quarter             236000 non-null  float64
14 Month               236000 non-null  object
15 Week of Year        236000 non-null  float64
16 Name of Day         236000 non-null  object
17 Hour                236000 non-null  float64
18 Event_Time          236000 non-null  object
dtypes: datetime64[ns](1), float64(11), object(7)
memory usage: 34.2+ MB

```

```
[5]: df.isnull().sum()
```

```

[5]: Index                1
Date                    1
event_type              1
product_id              1
category_id             1
category_code           1
brand                   1
price                   1
user_id                 1
user_session            1
State                  31371
User_Score              31371
Year                    1
Quarter                 1
Month                   1
Week of Year            1
Name of Day             1
Hour                    1
Event_Time              1
dtype: int64

```

```
[7]: df.skew()
```

```
[7]: Index          0.000000
     product_id    1.473684
     category_id    1.358517
     price          1.956892
     user_id        0.431609
     user_session   -0.000288
     User_Score     0.003175
     Year           0.000000
     Quarter        0.000000
     Week of Year   -0.421331
     Hour           -0.013025
     dtype: float64
```

```
[8]: df.kurtosis()
```

```
[8]: Index          -1.200000
     product_id     1.437978
     category_id     1.381260
     price          4.669445
     user_id        -0.981636
     user_session   -1.199220
     User_Score     -1.360660
     Year           0.000000
     Quarter        0.000000
     Week of Year   -1.017602
     Hour           -0.932056
     dtype: float64
```

```
[9]: column=['Index']
     df.drop(column, axis=1, inplace=True)
     df.head()
```

```
[9]:      Date event_type  product_id  category_id  \
0  2019-10-11      view      5016.0        183.0
1  2019-10-27      view      4240.0         43.0
2  2019-11-25      view      7721.0        208.0
3  2019-11-12      view       597.0         71.0
4  2019-11-13      view     13709.0         70.0

      category_code  brand  price  user_id  user_session  \
0  appliances.kitchen.oven  artel   36.01  1095536.0    5342766.0
1  electronics.video.tv    lg  2445.08  1153084.0    4600705.0
2  appliances.environment.vacuum  philips   257.38  364298.0    5991663.0
3  electronics.smartphone  huawei   163.20  3536496.0    9397681.0
4  electronics.telephone  nokia    21.85  3542250.0    8711820.0

      State  User_Score  Year  Quarter  Month  Week of Year  Name of Day  \
```

0	PA	3.0	2019.0	4.0	October	41.0	Friday
1	CT	4.0	2019.0	4.0	October	44.0	Sunday
2	HI	3.0	2019.0	4.0	November	48.0	Monday
3	NH	3.0	2019.0	4.0	November	46.0	Tuesday
4	NV	1.0	2019.0	4.0	November	46.0	Wednesday

	Hour	Event_Time
0	16.0	16:57:06
1	14.0	14:23:53
2	12.0	12:57:47
3	19.0	19:50:55
4	4.0	04:59:13

```
[13]: df.describe()
```

```
[13]:
```

	product_id	category_id	price	user_id \
count	236000.000000	236000.000000	236000.000000	2.360000e+05
mean	5211.558271	94.114797	390.030293	1.587240e+06
std	6334.376969	55.245882	382.386132	1.136710e+06
min	1.000000	0.000000	1.260000	9.900000e+01
25%	707.000000	71.000000	135.010000	5.734125e+05
50%	1197.000000	71.000000	249.660000	1.409442e+06
75%	7813.000000	115.000000	501.920000	2.461314e+06
max	26295.000000	305.000000	2574.040000	4.062342e+06

	user_session	User_Score	Year	Quarter	Week of Year \
count	2.360000e+05	204630.000000	236000.0	236000.0	236000.000000
mean	7.530784e+06	2.498045	2019.0	4.0	44.673119
std	4.347473e+06	1.118328	0.0	0.0	2.411971
min	1.500000e+01	1.000000	2019.0	4.0	40.000000
25%	3.758475e+06	1.000000	2019.0	4.0	43.000000
50%	7.543186e+06	2.000000	2019.0	4.0	45.000000
75%	1.129546e+07	3.000000	2019.0	4.0	47.000000
max	1.506575e+07	4.000000	2019.0	4.0	48.000000

	Hour
count	236000.000000
mean	11.091169
std	5.225956
min	0.000000
25%	7.000000
50%	11.000000
75%	15.000000
max	23.000000

```
[14]: import datetime as dt
NOW = dt.date(2019,12,30)
```

```
[15]: df['date'] = pd.DatetimeIndex(df.Date).date
```

```
[16]: df.head()
```

```
[16]:
```

	Date	event_type	product_id	category_id	\
0	2019-10-11	view	5016.0	183.0	
1	2019-10-27	view	4240.0	43.0	
2	2019-11-25	view	7721.0	208.0	
3	2019-11-12	view	597.0	71.0	
4	2019-11-13	view	13709.0	70.0	

	category_code	brand	price	user_id	user_session	\
0	appliances.kitchen.oven	artel	36.01	1095536.0	5342766.0	
1	electronics.video.tv	lg	2445.08	1153084.0	4600705.0	
2	appliances.environment.vacuum	philips	257.38	364298.0	5991663.0	
3	electronics.smartphone	huawei	163.20	3536496.0	9397681.0	
4	electronics.telephone	nokia	21.85	3542250.0	8711820.0	

	State	User_Score	Year	Quarter	Month	Week of Year	Name of Day	\
0	PA	3.0	2019.0	4.0	October	41.0	Friday	
1	CT	4.0	2019.0	4.0	October	44.0	Sunday	
2	HI	3.0	2019.0	4.0	November	48.0	Monday	
3	NH	3.0	2019.0	4.0	November	46.0	Tuesday	
4	NV	1.0	2019.0	4.0	November	46.0	Wednesday	

	Hour	Event_Time	date
0	16.0	16:57:06	2019-10-11
1	14.0	14:23:53	2019-10-27
2	12.0	12:57:47	2019-11-25
3	19.0	19:50:55	2019-11-12
4	4.0	04:59:13	2019-11-13

```
[32]: df_recency = df.groupby(['user_id'],as_index=False)['date'].max()
df_recency.columns = ['user_id','Last_Purchase_Date']
```

```
[33]: df_recency['Recency'] = df_recency.Last_Purchase_Date.apply(lambda x:(NOW - x).
↳days)
```

```
[34]: df_recency.head()
```

```
[34]:
```

	user_id	Last_Purchase_Date	Recency
0	99.0	2019-11-12	48
1	111.0	2019-11-30	30
2	226.0	2019-11-28	32
3	239.0	2019-11-13	47
4	244.0	2019-11-19	41

```
[35]: df_recency.drop(columns=['Last_Purchase_Date'],inplace=True)
```

```
[36]: FM_Table = df.groupby('user_id').agg({'product_id' : lambda x:len(x),
                                           'price' : lambda x:x.sum()})
```

```
[38]: FM_Table.rename(columns = {'product_id' : 'Frequency',
                                'price': 'Monetary_Value'},inplace= True)
```

```
[39]: FM_Table.head()
```

```
[39]:
```

	Frequency	Monetary_Value
user_id		
99.0	1.0	257.15
111.0	1.0	257.09
226.0	1.0	743.62
239.0	1.0	360.09
244.0	1.0	97.56

```
[40]: RFM_Table = df_recency.merge(FM_Table,left_on='user_id',right_on='user_id')
RFM_Table.head()
```

```
[40]:
```

	user_id	Recency	Frequency	Monetary_Value
0	99.0	48	1.0	257.15
1	111.0	30	1.0	257.09
2	226.0	32	1.0	743.62
3	239.0	47	1.0	360.09
4	244.0	41	1.0	97.56

```
[42]: df[df.user_id == 99]
```

```
[42]:
```

	Date	event_type	product_id	category_id	category_code	\
66708	2019-11-12	view	16570.0	91.0	furniture.bedroom.bed	

	brand	price	user_id	user_session	State	User_Score	Year	Quarter	\
66708	sv	257.15	99.0	11554527.0	CO	4.0	2019.0	4.0	

	Month	Week of Year	Name of Day	Hour	Event_Time	date
66708	November	46.0	Tuesday	3.0	03:30:32	2019-11-12

```
[44]: (NOW - dt.date(2019,11,12)).days == 48
```

```
[44]: True
```

```
[45]: quantiles = RFM_Table.quantile(q=[0.25,0.50,0.75])
quantiles = quantiles.to_dict()
```

```
[46]: segmented_rfm = RFM_Table.copy()
```



```
[48]: def RScore(x,p,d):
        if x <= d[p][0.25]:
            return 1
        elif x <= d[p][0.50]:
            return 2
        elif x <= d[p][0.75]:
            return 3
        else:
            return 4

    def FMScore(x,p,d):
        if x <= d[p][0.25]:
            return 4
        elif x <= d[p][0.50]:
            return 3
        elif x <= d[p][0.75]:
            return 2
        else:
            return 1
```

```
[49]: segmented_rfm['R_quartile'] = segmented_rfm['Recency'].apply(RScore,
    ↪args=('Recency',quantiles))
segmented_rfm['F_quartile'] = segmented_rfm['Frequency'].apply(FMScore,
    ↪args=('Frequency',quantiles))
segmented_rfm['M_quartile'] = segmented_rfm['Monetary_Value'].apply(FMScore,
    ↪args=('Monetary_Value',quantiles))
segmented_rfm.head()
```

```
[49]:
```

	user_id	Recency	Frequency	Monetary_Value	R_quartile	F_quartile	\
0	99.0	48	1.0	257.15	2	4	
1	111.0	30	1.0	257.09	1	4	
2	226.0	32	1.0	743.62	1	4	
3	239.0	47	1.0	360.09	2	4	
4	244.0	41	1.0	97.56	1	4	

	M_quartile
0	3
1	3
2	1
3	2
4	4

```
[50]: segmented_rfm['RFM_Segment'] = segmented_rfm.R_quartile.map(str)+segmented_rfm.
    ↪F_quartile.map(str)+segmented_rfm.M_quartile.map(str)
segmented_rfm.head()
```

```
[50]:
```

	user_id	Recency	Frequency	Monetary_Value	R_quartile	F_quartile	\
0	99.0	48	1.0	257.15	2	4	
1	111.0	30	1.0	257.09	1	4	
2	226.0	32	1.0	743.62	1	4	
3	239.0	47	1.0	360.09	2	4	
4	244.0	41	1.0	97.56	1	4	

	M_quartile	RFM_Segment
0	3	243
1	3	143
2	1	141
3	2	242
4	4	144

```
[51]: segmented_rfm['RFM_Score'] =
↳ segmented_rfm[['R_quartile', 'F_quartile', 'M_quartile']].sum(axis=1)
segmented_rfm.head()
```

```
[51]:
```

	user_id	Recency	Frequency	Monetary_Value	R_quartile	F_quartile	\
0	99.0	48	1.0	257.15	2	4	
1	111.0	30	1.0	257.09	1	4	
2	226.0	32	1.0	743.62	1	4	
3	239.0	47	1.0	360.09	2	4	
4	244.0	41	1.0	97.56	1	4	

	M_quartile	RFM_Segment	RFM_Score
0	3	243	9
1	3	143	8
2	1	141	6
3	2	242	8
4	4	144	9

```
[52]: print("Best Customers:␣
↳", len(segmented_rfm[segmented_rfm['RFM_Segment']=='111']))
print('Loyal Customers: ', len(segmented_rfm[segmented_rfm['F_quartile']==1]))
print("Big Spenders: ", len(segmented_rfm[segmented_rfm['M_quartile']==1]))
print('Almost Lost: ', len(segmented_rfm[segmented_rfm['RFM_Segment']=='134']))
print('Lost Customers:␣
↳', len(segmented_rfm[segmented_rfm['RFM_Segment']=='344']))
print('Lost Cheap Customers:␣
↳', len(segmented_rfm[segmented_rfm['RFM_Segment']=='444']))
```

```
Best Customers: 5089
Loyal Customers: 24182
Big Spenders: 51153
Almost Lost: 0
Lost Customers: 12178
Lost Cheap Customers: 12928
```

```
[53]: segmented_rfm['RFM_Score'].unique()
```

```
[53]: array([ 9,  8,  6,  7, 12, 10,  3, 11,  5,  4], dtype=int64)
```

```
[54]: segmented_rfm.groupby('RFM_Score').agg({
      'Recency': 'mean',
      'Frequency': 'mean',
      'Monetary_Value': ['mean', 'count'] }).round(1)
```

```
[54]:
```

	Recency	Frequency	Monetary_Value	
	mean	mean	mean	count
RFM_Score				
3	37.5	2.6	1450.8	5089
4	42.5	2.3	981.6	6651
5	51.5	2.2	829.5	5650
6	45.3	1.3	923.3	14110
7	43.8	1.1	643.4	22048
8	47.3	1.0	487.1	33055
9	54.5	1.0	384.6	46111
10	62.2	1.0	225.8	34179
11	70.5	1.0	145.2	24809
12	80.2	1.0	83.5	12928

```
[77]: RFM_Table.describe()
```

```
[77]:
```

	user_id	Recency	Frequency	Monetary_Value
count	2.046300e+05	204630.00000	204630.000000	204630.000000
mean	1.618077e+06	55.50735	1.153301	449.822358
std	1.142119e+06	16.82879	0.505220	492.311531
min	9.900000e+01	30.00000	1.000000	1.260000
25%	6.026678e+05	43.00000	1.000000	143.180000
50%	1.444042e+06	51.00000	1.000000	271.460000
75%	2.512144e+06	70.00000	1.000000	579.160000
max	4.062342e+06	90.00000	30.000000	13589.550000